



NVIDIA CUDA GETTING STARTED GUIDE FOR MAC OS X

DU-05348-001_v6.5 | August 2014

Installation and Verification on Mac OS X



TABLE OF CONTENTS

| | |
|--|-----------|
| Chapter 1. Introduction..... | 1 |
| 1.1. System Requirements..... | 1 |
| 1.2. About This Document..... | 2 |
| Chapter 2. Prerequisites..... | 3 |
| 2.1. CUDA-capable GPU..... | 3 |
| 2.2. Mac OS X Version..... | 3 |
| 2.3. Command-Line Tools..... | 3 |
| Chapter 3. Installation..... | 5 |
| 3.1. Download..... | 5 |
| 3.2. Install..... | 5 |
| 3.3. Uninstall..... | 6 |
| Chapter 4. Verification..... | 7 |
| 4.1. Driver..... | 7 |
| 4.2. Compiler..... | 7 |
| 4.3. Runtime..... | 8 |
| Chapter 5. Additional Considerations..... | 10 |

Chapter 1.

INTRODUCTION

CUDA[®] is a parallel computing platform and programming model invented by NVIDIA. It enables dramatic increases in computing performance by harnessing the power of the graphics processing unit (GPU).

CUDA was developed with several design goals in mind:

- ▶ Provide a small set of extensions to standard programming languages, like C, that enable a straightforward implementation of parallel algorithms. With CUDA C/C++, programmers can focus on the task of parallelization of the algorithms rather than spending time on their implementation.
- ▶ Support heterogeneous computation where applications use both the CPU and GPU. Serial portions of applications are run on the CPU, and parallel portions are offloaded to the GPU. As such, CUDA can be incrementally applied to existing applications. The CPU and GPU are treated as separate devices that have their own memory spaces. This configuration also allows simultaneous computation on the CPU and GPU without contention for memory resources.

CUDA-capable GPUs have hundreds of cores that can collectively run thousands of computing threads. These cores have shared resources including a register file and a shared memory. The on-chip shared memory allows parallel tasks running on these cores to share data without sending it over the system memory bus.

This guide will show you how to install and check the correct operation of the CUDA development tools.

1.1. System Requirements

To use CUDA on your system, you need to have:

- ▶ a CUDA-capable GPU
- ▶ Mac OS X 10.8 or later
- ▶ the *gcc* or *Clang* compiler and toolchain installed using Xcode
- ▶ the NVIDIA CUDA Toolkit (available from the [CUDA Download page](#))

Table 1 Mac Operating System Support in CUDA 6.5

| Operating System | Native x86_64 | GCC | Clang |
|------------------|---------------|-------|----------|
| Mac OS X 10.9.x | YES | | 5.0, 4.2 |
| Mac OS X 10.8.x | YES | 4.2.1 | 5.0 |

Before installing the CUDA Toolkit, you should read the [Release Notes](#), as they provide important details on installation and software functionality.

1.2. About This Document

This document is intended for readers familiar with the Mac OS X environment and the compilation of C programs from the command line. You do not need previous experience with CUDA or experience with parallel computation.

Chapter 2.

PREREQUISITES

2.1. CUDA-capable GPU

To verify that your system is CUDA-capable, under the **Apple** menu select **About This Mac**, click the **More Info ...** button, and then select **Graphics/Displays** under the **Hardware** list. There you will find the vendor name and model of your graphics card. If it is an NVIDIA card that is listed on the [CUDA-supported GPUs](#) page, your GPU is CUDA-capable.

The [Release Notes](#) for the CUDA Toolkit also contain a list of supported products.

2.2. Mac OS X Version

The CUDA Development Tools require an Intel-based Mac running Mac OSX v. 10.8 or later. To check which version you have, go to the **Apple** menu on the desktop and select **About This Mac**.

2.3. Command-Line Tools

The CUDA Toolkit requires that the native command-line tools (gcc, clang,...) are already installed on the system.

To install those command-line tools, Xcode must be installed first. Xcode is available from the Mac App Store.

Once Xcode is installed, the command-line tools can be installed by launching Xcode and following those steps:

1. **Xcode > Preferences... > Downloads > Components**
2. Install the *Command Line Tools* package

Alternatively, you can install the command-line tools from the Terminal window by typing the following command: `xcode-select --install`.

You can verify that the toolchain is installed by entering the command `/usr/bin/cc --help` from a **Terminal** window.

Chapter 3.

INSTALLATION

3.1. Download

Once you have verified that you have a supported NVIDIA GPU, a supported version of the MAC OS, and gcc, you need to download the NVIDIA CUDA Toolkit.

The NVIDIA CUDA Toolkit is available at no cost from the main [CUDA Downloads](#) page. It contains the driver and tools needed to create, build and run a CUDA application as well as libraries, header files, CUDA samples source code, and other resources.

The download can be verified by comparing the [posted MD5 checksum](#) with that of the downloaded file. If either of the checksums differ, the downloaded file is corrupt and needs to be downloaded again.

To calculate the MD5 checksum of the downloaded file, run the following:

```
$ openssl md5 <file>
```

3.2. Install

Use the following procedure to successfully install the CUDA driver and the CUDA toolkit. The CUDA driver and the CUDA toolkit must be installed for CUDA to function. If you have not installed a stand-alone driver, install the driver provided with the CUDA Toolkit.

Choose which packages you wish to install. The packages are:

- ▶ **CUDA Driver:** This will install `/Library/Frameworks/CUDA.framework` and the UNIX-compatibility stub `/usr/local/cuda/lib/libcuda.dylib` that refers to it.
- ▶ **CUDA Toolkit:** The CUDA Toolkit supplements the CUDA Driver with compilers and additional libraries and header files that are installed into `/Developer/NVIDIA/CUDA-6.5` by default. Symlinks are created in `/usr/local/cuda/` pointing to their respective files in `/Developer/NVIDIA/CUDA-6.5/`. Previous

installations of the toolkit will be moved to `/Developer/NVIDIA/CUDA-#. #` to better support side-by-side installations.

- **CUDA Samples (read-only):** A read-only copy of the CUDA Samples is installed in `/Developer/NVIDIA/CUDA-6.5/samples`. Previous installations of the samples will be moved to `/Developer/NVIDIA/CUDA-#. #/samples` to better support side-by-side installations.

Set up the required environment variables:

```
export PATH=/Developer/NVIDIA/CUDA-6.5/bin:$PATH
export DYLD_LIBRARY_PATH=/Developer/NVIDIA/CUDA-6.5/lib:$DYLD_LIBRARY_PATH
```

In order to modify, compile, and run the samples, the samples must also be installed with write permissions. A convenience installation script is provided: **cuda-install-samples-6.5.sh**. This script is installed with the cuda-samples-6-5 package.



To run CUDA applications in console mode on MacBook Pro with both an integrated GPU and a discrete GPU, use the following settings before dropping to console mode:

1. Uncheck **System Preferences > Energy Saver > Automatic Graphic Switch**
2. Drag the *Computer sleep* bar to *Never* in **System Preferences > Energy Saver**

3.3. Uninstall

The CUDA Driver, Toolkit and Samples can be uninstalled by executing the uninstall script provided with the Toolkit:

```
/Developer/NVIDIA/CUDA-6.5/bin/uninstall
```


Chapter 4.

VERIFICATION

Before continuing, it is important to verify that the CUDA toolkit can find and communicate correctly with the CUDA-capable hardware. To do this, you need to compile and run some of the included sample programs.



Ensure the `PATH` and `DYLD_LIBRARY_PATH` variables are **set correctly**.

4.1. Driver

If the CUDA Driver is installed correctly, the CUDA kernel extension (`/System/Library/Extensions/CUDA.kext`) should be loaded automatically at boot time. To verify that it is loaded, use the command

```
kextstat | grep -i cuda
```

4.2. Compiler

The installation of the compiler is first checked by running `nvcc -V` in a terminal window. The `nvcc` command runs the compiler driver that compiles CUDA programs. It calls the host compiler for C code and the NVIDIA PTX compiler for the CUDA code.



On Mac OS 10.8 with XCode 5, `nvcc` must be invoked with `--ccbin=path-to-clang-executable`. There are some features that are not yet supported: Clang language extensions (see <http://clang.llvm.org/docs/LanguageExtensions.html>), LLVM libc++ (only GNU libstdc++ is currently supported), language features introduced in C++11, and the `__global__` function template explicit instantiation definition.

The NVIDIA CUDA Toolkit includes CUDA sample programs in source form. To fully verify that the compiler works properly, a couple of samples should be built. After switching to the directory where the samples were installed, type:

```
make -C 0_Simple/vectorAdd
make -C 0_Simple/vectorAddDrv
make -C 1_Uutilities/deviceQuery
make -C 1_Uutilities/bandwidthTest
```

The builds should produce no error message. The resulting binaries will appear under **<dir>/bin/x86_64/darwin/release**. To go further and build all the CUDA samples, simply type **make** from the samples root directory.

4.3. Runtime

After compilation, go to **bin/x86_64/darwin/release** and run **deviceQuery**. If the CUDA software is installed and configured correctly, the output for **deviceQuery** should look similar to that shown in [Figure 1](#).

```
./deviceQuery Starting...

CUDA Device Query (Runtime API) version (CUDART static linking)

Detected 1 CUDA Capable device(s)

Device 0: "GeForce GT 650M"
  CUDA Driver Version / Runtime Version      6.0 / 6.0
  CUDA Capability Major/Minor version number: 3.0
  Total amount of global memory:              512 MBytes (536543232 bytes)
  ( 2) Multiprocessors, (192) CUDA Cores/MP: 384 CUDA Cores
  GPU Clock rate:                            405 MHz (0.41 GHz)
  Memory Clock rate:                         2000 Mhz
  Memory Bus Width:                          128-bit
  L2 Cache Size:                             262144 bytes
  Maximum Texture Dimension Size (x,y,z)      1D=(65536), 2D=(65536, 65536), 3D=(4096, 4096, 4096)
  Maximum Layered 1D Texture Size, (num) layers 1D=(16384), 2048 layers
  Maximum Layered 2D Texture Size, (num) layers 2D=(16384, 16384), 2048 layers
  Total amount of constant memory:             65536 bytes
  Total amount of shared memory per block:     49152 bytes
  Total number of registers available per block: 65536
  Warp size:                                  32
  Maximum number of threads per multiprocessor: 2048
  Maximum number of threads per block:         1024
  Max dimension size of a thread block (x,y,z): (1024, 1024, 64)
  Max dimension size of a grid size    (x,y,z): (2147483647, 65535, 65535)
  Maximum memory pitch:                       2147483647 bytes
  Texture alignment:                          512 bytes
  Concurrent copy and kernel execution:        Yes with 1 copy engine(s)
  Run time limit on kernels:                   Yes
  Integrated GPU sharing Host Memory:           No
  Support host page-locked memory mapping:      Yes
  Alignment requirement for Surfaces:           Yes
  Device has ECC support:                      Disabled
  Device supports Unified Addressing (UVA):      Yes
  Device PCI Bus ID / PCI location ID:          1 / 0
  Compute Mode:
    < Default (multiple host threads can use ::cudaSetDevice() with device simultaneously) >

deviceQuery, CUDA Driver = CUDART, CUDA Driver Version = 6.0, CUDA Runtime Version = 6.0, NumDevs = 1, Device0 = GeForce GT 650M
Result = PASS
```

Figure 1 Valid Results from deviceQuery CUDA Sample

Note that the parameters for your CUDA device will vary. The key lines are the first and second ones that confirm a device was found and what model it is. Also, the next-to-last line, as indicated, should show that the test passed.

Running the **bandwidthTest** sample ensures that the system and the CUDA-capable device are able to communicate correctly. Its output is shown in [Figure 2](#)

```

[CUDA Bandwidth Test] - Starting...
Running on...

Device 0: Quadro FX 4800
Quick Mode

Host to Device Bandwidth, 1 Device(s)
PINNED Memory Transfers
  Transfer Size (Bytes)      Bandwidth(MB/s)
  33554432                  5287.2

Device to Host Bandwidth, 1 Device(s)
PINNED Memory Transfers
  Transfer Size (Bytes)      Bandwidth(MB/s)
  33554432                  3901.9

Device to Device Bandwidth, 1 Device(s)
PINNED Memory Transfers
  Transfer Size (Bytes)      Bandwidth(MB/s)
  33554432                  37519.2

Result = PASS
dhcp-172-16-194-251:release eyoung$

```

Figure 2 Valid Results from bandwidthTest CUDA Sample

Note that the measurements for your CUDA-capable device description will vary from system to system. The important point is that you obtain measurements, and that the second-to-last line (in [Figure 2](#)) confirms that all necessary tests passed.

Should the tests not pass, make sure you have a CUDA-capable NVIDIA GPU on your system and make sure it is properly installed.

If you run into difficulties with the link step (such as libraries not being found), consult the *Release Notes* found in the `doc` folder in the CUDA Samples directory.

To see a graphical representation of what CUDA can do, run the particles executable.

Chapter 5.

ADDITIONAL CONSIDERATIONS

Now that you have CUDA-capable hardware and the NVIDIA CUDA Toolkit installed, you can examine and enjoy the numerous included programs. To begin using CUDA to accelerate the performance of your own applications, consult the [CUDA C Programming Guide](#).

A number of helpful development tools are included in the CUDA Toolkit to assist you as you develop your CUDA programs, such as NVIDIA® Nsight™ Eclipse Edition, NVIDIA Visual Profiler, `cuda-gdb`, and `cuda-memcheck`.

For technical support on programming questions, consult and participate in the [Developer Forums](#).

Notice

ALL NVIDIA DESIGN SPECIFICATIONS, REFERENCE BOARDS, FILES, DRAWINGS, DIAGNOSTICS, LISTS, AND OTHER DOCUMENTS (TOGETHER AND SEPARATELY, "MATERIALS") ARE BEING PROVIDED "AS IS." NVIDIA MAKES NO WARRANTIES, EXPRESSED, IMPLIED, STATUTORY, OR OTHERWISE WITH RESPECT TO THE MATERIALS, AND EXPRESSLY DISCLAIMS ALL IMPLIED WARRANTIES OF NONINFRINGEMENT, MERCHANTABILITY, AND FITNESS FOR A PARTICULAR PURPOSE.

Information furnished is believed to be accurate and reliable. However, NVIDIA Corporation assumes no responsibility for the consequences of use of such information or for any infringement of patents or other rights of third parties that may result from its use. No license is granted by implication of otherwise under any patent rights of NVIDIA Corporation. Specifications mentioned in this publication are subject to change without notice. This publication supersedes and replaces all other information previously supplied. NVIDIA Corporation products are not authorized as critical components in life support devices or systems without express written approval of NVIDIA Corporation.

Trademarks

NVIDIA and the NVIDIA logo are trademarks or registered trademarks of NVIDIA Corporation in the U.S. and other countries. Other company and product names may be trademarks of the respective companies with which they are associated.

Copyright

© 2009-2014 NVIDIA Corporation. All rights reserved.